



COMPUTATIONAL FINANCE & RISK MANAGEMENT

UNIVERSITY *of* WASHINGTON

Department of Applied Mathematics

R Introduction

Guy Yollin

Acting Instructor, Applied Mathematics
University of Washington

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

Lecture references

-  W. N. Venables and D. M. Smith.
An Introduction to R.
2011.

-  J. Adler.
R in a Nutshell: A Desktop Quick Reference.
O'Reilly Media, 2010.

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3


- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

What is R?

- R is a *language* and *environment* for statistical computing and graphics
- R is based on the *S language* originally developed by John Chambers and colleagues at AT&T Bell Labs in the late 1970s and early 1980s
- R (sometimes called “*GNU S*”) is free open source software licensed under the GNU general public license (GPL 2)
- R development was initiated by Robert Gentleman and Ross Ihaka at the University of Auckland, New Zealand
- R is formally known as The R Project for Statistical Computing
 - www.r-project.org

What is R great at?

- Data manipulation
- Data analysis
- Statistical modeling
- Data visualization



Plot from the `PerformanceAnalytics` package

S language implementations

R is the most recent and full-featured implementation of the S language

- Original S - AT & T Bell Labs
- S-PLUS (S plus a GUI)
 - Statistical Sciences, Inc.[†]
 - Mathsoft, Inc.
 - Insightful, Inc.
 - Tibco, Inc.
- R - The R Project for Statistical Computing





Figure from *The History of S and R*, John Chambers, 2006

[†]Founded by UW Professor Doug Martin, CompFin Program Director

R timeline

The R Timeline



Recognition of software excellence

Association for Computing Machinery

John Chambers received the 1998 ACM Software System Award

Dr. Chambers' work will forever alter the way people analyze, visualize, and manipulate data

American Statistical Association

Robert Gentleman and Ross Ihaka received the 2009 ASA Statistical Computing and Graphics Award

In recognition for their work in initiating the R Project for Statistical Computing

The R Foundation

The R Foundation is the non-profit organization located in Vienna, Austria which is responsible for developing and maintaining R

- Hold and administer the copyright of R software and documentation
- Support continued development of R
- Organize meetings and conferences related to statistical computing
- Officers

Presidents Robert Gentleman, Ross Ihaka

Secretary Friedrich Leisch

Treasurer Kurt Hornik

At Large John Chambers

Auditors Peter Dalgaard, Martin Maechler

The R Core Team

- Douglas Bates – University of Wisconsin Madison
- John Chambers – Stanford University
- Peter Dalgaard – University of Copenhagen
- Seth Falcon – Fred Hutchinson Cancer Research Center
- Robert Gentleman – Genetech
- Kurt Hornik – Vienna University of Economics and Business
- Stefano Iacus – University of Milan
- Ross Ihaka – University of Auckland
- Friedrich Leisch – Ludwig-Maximilians –University Munich
- Uwe Ligges – TU Dortmund University
- Thomas Lumley – University of Auckland
- Martin Maechler – ETH Swiss Federal Institute of Technology Zurich
- Duncan Murdoch – University of Western Ontario
- Paul Murrell – University of Auckland
- Martyn Plummer – International Agency for Research on Cancer
- Brian Ripley – University of Oxford
- Deepayan Sarkar – Fred Hutchinson Cancer Research Center
- Duncan Temple Lang – University of California Davis
- Luke Tierney – University of Iowa
- Simon Urbanek – AT & T Research Labs

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2



- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R



R language references

- R in a Nutshell: A Desktop Quick Reference
 - Joseph Adler
 - O'Reilly Media, 2009
- The Art of R Programming
 - Norman Matloff
 - No Starch Press, 2011



Other introductory R texts

- **Introductory Statistics with R**
2nd Edition
 - P. Dalgaard
 - Springer, 2008
- **A Beginner's Guide to R**
 - Zuur, Ieno, Meesters
 - Springer, 2009



Experience with other statistical computing languages

For those with experience in MATLAB, David Hiebeler has created a MATLAB/R cross reference document:

- <http://www.math.umaine.edu/~hiebeler/comp/matlabR.pdf>

For those with experience in SAS, SPSS, or Stata, Robert Muenchen has written R books for this audience:

- <http://r4stats.com>

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- **R language and environment basics**
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R


The R GUI

Running R in Windows

- Typically run Rgui.exe



- Can also run R.exe from command prompt
- Or run Rterm.exe in batch mode



The R GUI on a Windows platform

Interactive R session

- R is an *interpreted* language
- The R GUI is an *interactive command driven environment*
 - type R commands at the R GUI console
 - Run previously created R scripts (R commands in a text file)



Commands entered interactively into the R console

Assigning values to variables

- Typical variable assignment
 - assignment operator: <-
 - assignment function: assign
 - equal sign: =
 - must be used to assign arguments in a function call
- Special purpose assignment
 - global assignment operator: <<-
- Deprecated assignment operator
 - underscore character: _

R Code: Variable assignment

```
> y <- 5
> y
[1] 5

> assign("e",2.7183)
> e

[1] 2.7183

> s = sqrt(2)
> s
[1] 1.414214

> r <- rnorm(n=2)
> r
[1] 0.1323967 0.7617530
```

Object orientation in R

Everything in R is an Object

- Use functions `ls` and `objects` to list all objects in the current workspace

R Code: Listing objects

```
> x <- c(3.1416,2.7183)
> m <- matrix(rnorm(9),nrow=3)
> tab <- data.frame(store=c("downtown","eastside","airport"),sales=c(32,17,24))
> cities <- c("Seattle","Portland","San Francisco")
> ls()
[1] "cities"      "e"           "filename"    "m"           "r"           "s"           "tab"
[8] "x"           "y"
```

Data types

All R objects have a *type* or *storage mode*

- Use function `typeof` to display an object's type
- Common types are:
 - double
 - character
 - list
 - integer

R Code: Object type (storage mode)

```
> x  
[1] 3.1416 2.7183  
> typeof(x)  
[1] "double"  
  
> cities  
[1] "Seattle"      "Portland"  
[3] "San Francisco"  
> typeof(cities)  
[1] "character"
```

Object classes

All R objects have a *class*

- Use function `class` to display an object's class
- There are many R classes; basic classes are:
 - numeric
 - character
 - `data.frame`
 - `matrix`

R Code: Object class

```
> m
      [,1]      [,2]      [,3]
[1,] 1.5250413  0.6412772  2.52021809
[2,] 1.6394314 -0.6197558  0.80289579
[3,] 0.8032637  0.2847256 -0.03179198

> class(m)
[1] "matrix"

> tab
      store sales
1 downtown    32
2 eastside     17
3 airport      24

> class(tab)
[1] "data.frame"
```

Vectors

R is a vector/matrix language

- vectors can easily be created with `c`, the combine function
- most places where single value can be supplied, a vector can be supplied and R will perform a vectorized operation

R Code: Creating vectors and vector operations

```
> constants <- c(3.1416,2.7183,1.4142,1.6180)
> names(constants) <- c("pi","euler","sqrt2","golden")
> constants

    pi   euler   sqrt2   golden
3.1416 2.7183 1.4142 1.6180

> constants^2

    pi     euler     sqrt2     golden
9.869651 7.389155 1.999962 2.617924

> 10*constants

    pi   euler   sqrt2   golden
31.416 27.183 14.142 16.180
```

Indexing vectors

Vectors indices are placed with square brackets: []

Vectors can be indexed in any of the following ways:

- vector of positive integers
- vector of negative integers
- vector of named items
- logical vector

R Code: Indexing vectors

```
> constants[c(1,3,4)]  
pi  sqrt2 golden  
3.1416 1.4142 1.6180  
  
> constants[c(-1,-2)]  
sqrt2 golden  
1.4142 1.6180  
  
> constants[c("pi","golden")]  
pi golden  
3.1416 1.6180  
  
> constants > 2  
pi euler sqrt2 golden  
TRUE TRUE FALSE FALSE  
  
> constants[constants > 2]  
pi euler  
3.1416 2.7183
```

The recycling rule

When 2 vectors of unequal length are involved in an operation, the shorter one is recycled to equal the length of the longer vector

R Code: Illustration of recycling

```
> constants  
  
pi euler sqrt2 golden  
3.1416 2.7183 1.4142 1.6180  
  
> constants*2  
  
pi euler sqrt2 golden  
6.2832 5.4366 2.8284 3.2360  
  
> constants*c(0,1)  
  
pi euler sqrt2 golden  
0.0000 2.7183 0.0000 1.6180  
  
> constants*c(0,1,2)  
  
pi euler sqrt2 golden  
0.0000 2.7183 2.8284 0.0000
```

last input generates a warning: longer object length
is not a multiple of shorter object length

Sequences

An integer sequence vector can be created with the `:` operator

A general numeric sequence vector can be created with the `seq` function

R Code: seq arguments

```
> args(seq.default)  
  
function (from = 1, to = 1, by = ((to - from)/(length.out - 1)),  
        length.out = NULL, along.with = NULL, ...)  
NULL
```

`from` starting value

`to` ending value

`by` increment

`len` length of sequence

Sequences

R Code: Creating sequences

```
> 1:5  
[1] 1 2 3 4 5  
  
> -5:5  
[1] -5 -4 -3 -2 -1  0  1  2  3  4  5  
  
> seq(from=0,to=1,len=5)  
[1] 0.00 0.25 0.50 0.75 1.00  
  
> seq(from=0,to=20,by=2.5)  
[1]  0.0  2.5  5.0  7.5 10.0 12.5 15.0 17.5 20.0
```

Passing arguments to functions

- unnamed arguments are assigned according to their position
- named arguments are assigned according to their name and can be in any position
- partial name matching is performed
- arguments with default values are not required to be passed

R Code: Illustration of flexibility in passing arguments

```
> seq(0,10,2)
[1] 0 2 4 6 8 10

> seq(by=2,0,10)
[1] 0 2 4 6 8 10

> seq(0,10,len=5)
[1] 0.0 2.5 5.0 7.5 10.0

> seq(0,10)
[1] 0 1 2 3 4 5 6 7 8 9 10
```

The "..." argument

Many functions include in their argument list a ...

R Code: The plot function arguments

```
> args(plot.default)  
  
function (x, y = NULL, type = "p", xlim = NULL, ylim = NULL,  
        log = "", main = NULL, sub = NULL, xlab = NULL, ylab = NULL,  
        ann = par("ann"), axes = TRUE, frame.plot = axes, panel.first = NULL,  
        panel.last = NULL, asp = NA, ...)  
NULL
```

- This is a mechanism to allow additional arguments to be passed which will subsequently be passed on to a sub-function that the main function will call
- An example of this would be passing graphic parameters (e.g. `lwd=2`) to the `plot` function which will subsequently call and pass these arguments on to the `par` function

The rep function

The rep function is used to create (or initialize) vectors

R Code: Examples of rep

```
> rep(0,10)    # initialize a vector  
[1] 0 0 0 0 0 0 0 0 0 0  
  
> rep(1:4, 2) # repeat pattern 2 times  
[1] 1 2 3 4 1 2 3 4  
  
> rep(1:4, each = 2)  # repeat each element 2 times  
[1] 1 1 2 2 3 3 4 4  
  
> rep(1:4, c(2,1,2,1))  
[1] 1 1 2 3 3 4  
  
> rep(1:4, each = 2, len = 10)    # 8 integers plus two recycled 1's.  
[1] 1 1 2 2 3 3 4 4 1 1  
  
> rep(1:4, each = 2, times = 3)  # length 24, 3 complete replications  
[1] 1 1 2 2 3 3 4 4 1 1 2 2 3 3 4 4 1 1 2 2 3 3 4 4
```

Generic functions

A generic function behaves in a way that is appropriate based on the class of its argument; for example:

- plot
- print
- summary

R Code: Some classes handled by the plot function

```
> methods(plot)[1:15]
```

[1] "plot.acf"	"plot.data.frame"	"plot.decomposed.ts"
[4] "plot.default"	"plot.dendrogram"	"plot.density"
[7] "plot.ecdf"	"plot.factor"	"plot.formula"
[10] "plot.function"	"plot.hclust"	"plot.histogram"
[13] "plot.HoltWinters"	"plot.isoreg"	"plot.lm"

- generic functions implement simple polymorphism for S3 objects

R packages

- All R functions are stored in *packages*
- The standard R distribution includes *core* packages and *recommended* packages:
 - Core R packages
 - base, utils, stats, methods, graphics, grDevices, datasets
 - Recommended packages
 - boot, rpart, foreign, MASS, cluster, Matrix, etc.
 - Additional packages can be downloaded through the R GUI or via the `install.packages` function
- When R is initially loaded, only core R packages are loaded by default
 - Additional packages are loaded via the `library` command
 - Packages datasets are made accessible via the `data` command

Loading packages and data into your R session

The library and data functions are used to load additional libraries and data into the current R session

R Code: The library and data function

```
> args(library)

function (package, help, pos = 2, lib.loc = NULL, character.only = FALSE,
logical.return = FALSE, warn.conflicts = TRUE, quietly = FALSE,
keep.source = getOption("keep.source.pkgs"), verbose = getOption("verbose"))
NULL

> args(data)

function (... , list = character(), package = NULL, lib.loc = NULL,
verbose = getOption("verbose"), envir = .GlobalEnv)
NULL

> library(nutshell)
> data(top.bacon.searching.cities)
> top.bacon.searching.cities[1,]

  city rank
1 Seattle 100
```

Installing contributed packages

The `install.packages` function can be used to install contributed packages

R Code: The `install.packages` function


```
> args(install.packages)

function (pkgs, lib, repos = getOption("repos"), contriburl = contrib.url(repos,
  type), method, available = NULL, destdir = NULL, dependencies = NA,
  type = getOption("pkgType"), configure.args = getOption("configure.args"),
  configure.vars = getOption("configure.vars"), clean = FALSE,
  Ncpus = getOption("Ncpus", 1L), libs_only = FALSE, INSTALL_opts,
  ...)
NULL

> install.packages("nutshell")
> # or if repository needs to be specified
> install.packages("nutshell", repos="http://cran.fhcrc.org")
```

Installing contributed packages

Packages can also be installed through the R GUI



Packages for basic computational finance

The following R add-on packages are recommended for computational finance:

	Package	Description
PerformanceAnalytics	zoo	Time series objects
	tseries	Time series analysis and computational finance
	quantmod	Performance and risk analysis
	xts	Quantitative financial modeling framework
		Extensible time series

Writing R functions

One of the strengths of R is that it can easily be extended by writing new R functions; in fact, much of R is written in R

The syntax for defining a new R function is as follows:

```
name <- function(arg_1, arg_2, ...) expression
```

R Code: Create a user-defined function

```
> percentChange <- function(x)
{
  100*(x[-1]/x[-length(x)]-1)
}
> sales <- c(100,105,110,105,100)
> percentChange(sales)
[1]  5.000000  4.761905 -4.545455 -4.761905
```

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- **Data structures, data manipulation, working directory, data files**
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

The list object

A list object is a container that can hold other objects of different types

R Code: Creating lists

```
> myList <- list(pi=3.1416,euler=2.7183,golden=1.6180)
> class(myList)

[1] "list"

> length(myList)

[1] 3

> myList

$pi
[1] 3.1416

$euler
[1] 2.7183

$golden
[1] 1.618

> diverseList <- list(magic=myList,random=matrix(rnorm(4),ncol=2),
  state=c("WA","OR"))
```

Accessing items in a list

Items in a list can be accessed using `[]`, `[[]]`, or `$` syntax as follows:

- `[]` returns a sublist
 - vector of positive integers
 - vector of named items
 - logical vector
- `[[]]` returns a single element
 - single integer
 - single name
- `$` returns a single element
 - single name

R Code: Indexing lists

```
> myList[2]  
$euler  
[1] 2.7183  
  
> myList[[2]]  
[1] 2.7183  
  
> myList[["pi"]]  
[1] 3.1416  
  
> myList$golden  
[1] 1.618  
  
> diverseList[[3]][2]  
[1] "OR"
```

The data.frame object

- A data.frame is a 2D matrix-like object where the *columns* can be of *different* classes; for example:
 - column with dates
 - column with characters
 - column with integers
 - column with numeric

R Code: The data.frame object

```
> data(batting.2008)
> class(batting.2008)

[1] "data.frame"

> dim(batting.2008)

[1] 1384   32

> batting.2008[1:2,1:4]

  nameLast nameFirst weight height
1     Abreu      Bobby     200      72
2     Alou       Moises     190      75

> class(batting.2008[,2])

[1] "character"

> class(batting.2008[,3])

[1] "integer"

> class(batting.2008[,4])

[1] "numeric"
```

The head and tail functions

R Code: The head and tail functions

```
> args(getS3method("head","data.frame"))
```

```
function (x, n = 6L, ...)  
NULL
```

```
> data(dow30)  
> head(dow30)
```

	symbol	Date	Open	High	Low	Close	Volume	Adj.Close
1	MMM	2009-09-21	73.91	74.68	73.91	74.54	2560400	74.54
2	MMM	2009-09-18	75.12	75.25	74.50	74.62	4387900	74.62
3	MMM	2009-09-17	75.34	75.45	74.50	74.89	3371500	74.89
4	MMM	2009-09-16	74.76	75.49	74.50	75.38	2722500	75.38
5	MMM	2009-09-15	74.63	74.88	74.00	74.68	3566900	74.68
6	MMM	2009-09-14	73.72	74.64	73.42	74.56	3466400	74.56

```
> tail(dow30,3)
```

	symbol	Date	Open	High	Low	Close	Volume	Adj.Close
7480	DIS	2008-09-24	32.59	32.59	31.63	31.77	13600300	31.30
7481	DIS	2008-09-23	32.88	33.32	32.15	32.53	13450900	32.05
7482	DIS	2008-09-22	33.85	34.05	32.84	32.91	18394300	32.42

Size-related and diagnostic helper functions

R has a number of size related and diagnostic helper functions

Function	Description
dim	return dimensions of a multidimensional object
nrow	number of rows of a multidimensional object
ncol	number of columns of a multidimensional object
length	length a vector or list
head	display first n rows (elements)
tail	display last n rows (elements)
str	summarize structure of an object

Indexing data.frames and matrices

R has extremely powerful data manipulation capabilities especially in the area of vector and matrix indexing

- data.frames and matrices can be indexed in any of the following ways
 - vector of positive integers
 - vector of negative integers
 - character vector of columns (row) names
 - a logical vector
- Since data.frames are stored internally as lists, their columns can be accessed with the \$ operator as well

Indexing data.frames and matrices

R has extremely powerful data manipulation capabilities especially in the area of vector and matrix indexing

- data.frames and matrices can be indexed in any of the following ways
 - vector of positive integers
 - vector of negative integers
 - character vector of columns (row) names
 - a logical vector
- Since data.frames are stored internally as lists, their columns can be accessed with the \$ operator as well

R Code: Indexing 2D objects

```
> dow30[1:4,c("symbol","Date","Close")]
```

	symbol	Date	Close
1	MMM	2009-09-21	74.54
2	MMM	2009-09-18	74.62
3	MMM	2009-09-17	74.89
4	MMM	2009-09-16	75.38

```
> head(dow30[-1,c(1,2,6)],3)
```

	symbol	Date	Close
2	MMM	2009-09-18	74.62
3	MMM	2009-09-17	74.89
4	MMM	2009-09-16	75.38

```
> dow30[dow30[,"Volume"]>1.5e9,  
       c("symbol","Date","Close","Volume")]
```

	symbol	Date	Close	Volume
2047	C	2009-08-07	3.85	1898814600
2049	C	2009-08-05	3.58	2672492000
2159	C	2009-02-27	1.50	1868209400

Factors

A *factor* is a data type for representing categorial data

R Code: Create a factor variable

```
> pet.str <- c("dog", "cat", "cat", "dog", "fish", "dog", "rabbit")
> pets <- as.factor(pet.str)
> pets

[1] dog     cat     cat     dog     fish    dog     rabbit
Levels: cat dog fish rabbit

> as.numeric(pets)

[1] 2 1 1 2 3 2 4

> levels(pets)

[1] "cat"    "dog"    "fish"   "rabbit"
```

- factors are encoded as integers (starting at 1)
- the *levels* of a factor variable contain the categorical labels

The working directory

Unless overridden by a filename which includes a path, R reads and writes files to the *working directory*


R Code: Getting and setting the working directory

```
> (my.wd <- getwd())
[1] "C:/Rprojects/UW/RIntro"
> setwd(R.home())
> getwd()
[1] "C:/R/R-2.15.1"
> setwd(my.wd)
> getwd()
[1] "C:/Rprojects/UW/RIntro"
```

- The backslash character “\” in a character string is used to begin an escape sequence, so to use backslash in a string enter it as “\\”
- The forward slash character “/” can also be used as a directory separator on windows systems

The working directory

The working directory can also be changed from the R GUI



R Gui (64-bit)

File Edit View Misc Packages Windows Help

Source R code...
New script...
Open script...
Display file(s)...

Load Workspace...
Save Workspace...

Load History...
Save History...
Change dir...

Print...
Save to File...

Exit

```
3.14 2.72 1.41 1.62
s")= chr [1:4] "pi" "euler" "sqrt2" "golden"
3 obs. of  2 variables:
w/ 3 levels "airport","downtown",..: 2 3 1
2 17 24
nts)
   Min. 1st Qu. Median     Mean 3rd Qu.    Max.
1.414   1.567  2.168   2.223   2.824   3.142
> dim(constants)
NULL
> dim(tab)
[1] 3 2
> length(constants)
[1] 4
> length(tab)
[1] 2
> length(m)
[1] 9
> m
      [,1]      [,2]      [,3]
[1,] -0.4017547 -0.3146403 -0.04649634
[2,] -0.4864167  1.3425015  0.80295826
[3,] -0.1029465 -1.5566827 -0.92187629
```

The `read.table` function

The `read.table` function is used *extensively* to load data into R

R Code: `read.table` arguments

```
> args(read.table)

function (file, header = FALSE, sep = "", quote = "\"\"", dec = ".",
  row.names, col.names, as.is = !stringsAsFactors, na.strings = "NA",
  colClasses = NA, nrows = -1, skip = 0, check.names = TRUE,
  fill = !blank.lines.skip, strip.white = FALSE, blank.lines.skip = TRUE,
  comment.char = "#", allowEscapes = FALSE, flush = FALSE,
  stringsAsFactors = default.stringsAsFactors(), fileEncoding = "",
  encoding = "unknown", text)
NULL
```


`file` file name (with path if necessary)

`header` TRUE/FALSE if there are column names in the file

`sep` column separation character (e.g. comma or tab)

`as.is` tells R not to convert strings into factors

Reading a text file



The screenshot shows a Notepad++ window with the title bar "C:\Rprojects\UW\AMATH542\RIntro\INTC.CSV - Notepad++". The menu bar includes File, Edit, Search, View, Encoding, Language, Settings, Macro, Run, TextFX, Plugins, Window, and Help. The toolbar below has icons for file operations like Open, Save, Find, and Print. The main text area is titled "INTC.CSV" and contains 254 lines of data. The first few lines are:

```
1 Date,Open,High,Low,Close,Volume,Adj.Close
2 2009-09-21,19.47,19.76,19.42,19.54,36132900,19.54
3 2009-09-18,19.52,19.7,19.37,19.56,53482300,19.56
4 2009-09-17,19.63,19.89,19.22,19.41,64594100,19.41
5 2009-09-16,19.7,19.75,19.44,19.65,49427600,19.65
6 2009-09-15,19.54,19.71,19.36,19.55,48173000,19.55
7 2009-09-14,19.31,19.49,19.28,19.36,52357900,19.36
```

At the bottom, status bars show "Normal length : 12721 lines : 254", "Ln : 1 Col : 1 Sel : 0", "Dos\Windows", and "ANSI".

R Code: Read csv file

```
> dat <- read.table("intc.csv", header=TRUE, sep=",", as.is=TRUE)
> dat[1:5,]
```

	Date	Open	High	Low	Close	Volume	Adj.Close
1	2009-09-21	19.47	19.76	19.42	19.54	36132900	19.54
2	2009-09-18	19.52	19.70	19.37	19.56	53482300	19.56
3	2009-09-17	19.63	19.89	19.22	19.41	64594100	19.41
4	2009-09-16	19.70	19.75	19.44	19.65	49427600	19.65
5	2009-09-15	19.54	19.71	19.36	19.55	48173000	19.55

Writing text files

The functions `write.table` and `write` are used to write text files

R Code: `write.table` and `write` arguments

```
> args(write.table)

function (x, file = "", append = FALSE, quote = TRUE, sep = " ",
  eol = "\n", na = "NA", dec = ".", row.names = TRUE, col.names = TRUE,
  qmethod = c("escape", "double"), fileEncoding = "")

NULL

> args(write)

function (x, file = "data", ncolumns = if (is.character(x)) 1 else 5,
  append = FALSE, sep = " ")

NULL
```


- x** object to be written (data.frame, matrix, vector)
- file** file name (with path if necessary)
- sep** column separation character (e.g. comma or tab)
- row.names** write row names (T/F)
- col.names** write col names (T/F)

Writing text files

R Code: Write text files


```
> write(x=constants,file="vector.dat",sep="\t")
> write.table(x=m,file="matrix.dat",sep="\t",row.names=F,col.names=F)
> file.info(list.files(pattern=". [d] [a] [t]",full.names=T))[,c("size","mtime")]

      size          mtime
./matrix.dat  164 2012-08-31 08:15:15
./vector.dat   28 2012-08-31 08:15:15
```



A screenshot of Notepad++ showing the contents of the file 'matrix.dat'. The file contains a 4x3 matrix of numerical values:

	1	2	3
1	1.22063076123015	-0.883747093620611	-0.649465341856833
2	-2.19833622100779	2.0769886874053	-0.968967050221968
3	0.0266252084306489	1.02091620385764	-1.70296534035295
4			



A screenshot of Notepad++ showing the contents of the file 'vector.dat'. The file contains a 2x4 matrix of numerical values:

	1	2	3	4
1	3.1416	2.7183	1.4142	1.618
2				

- note the use of the `list.files` function, the `file.info` function, and the use of regular expressions

Functions to examine objects and their structures

These functions help to query and unpack an object

`class` query an objects class

`str` reports structure of an object

`attributes` returns list of objects attributes

`attr` get/set attributes of an object

`names` gets the names of a list, vector, data.frame, etc.

`dimnames` gets the row and column names of a data.frame or matrix

`colnames` column names of a data.frame or matrix

`rownames` row names of a data.frame or matrix

`dput` makes an ASCII representation of an object

`unclass` removes class attribute of an object

`unlist` converts a list to a vector

The paste function

The paste function concatenates (*pastes*) strings and numerical values together

- its like a flexible version of sprintf

R Code: The paste function

```
> args(paste)  
  
function (..., sep = " ", collapse = NULL)  
NULL  
  
> a <- 2; b <- 2  
> paste("We know that: ", a, " + ", b, " = ", a+b, sep = "")  
[1] "We know that: 2 + 2 = 4"  
  
> paste("variable", 1:5, sep = "")  
[1] "variable1" "variable2" "variable3" "variable4" "variable5"
```

The apply function

The apply function is an *extremely* useful function that *applies* a given function across the rows and/or columns of a matrix

R Code: The apply function

```
> args(apply)
function (X, MARGIN, FUN, ...)
NULL

> set.seed(1)
> (m <- matrix(sample(9), ncol=3))

 [,1] [,2] [,3]
[1,]    3    6    8
[2,]    9    2    7
[3,]    5    4    1

> apply(m, 2, sum)
[1] 17 12 16
```

- There are a number of *apply related* functions; one mark of mastering R is mastering apply related functions

S4 Classes

S4 classes are a more modern implementation of object-oriented programming in R compared to S3 classes

- Data in an S4 class is organized into *slots*; slots can be accessed using:
 - the @ operator: `object@name`
 - the slot function: `slot(object, name)`
- Methods for an S4 class can be queried with the `showMethods` function
 - `showMethods(class = "fGARCH")`
- Methods can be retrieved/viewed with the `getMethod` function
 - `getMethod("predict","fGARCH")`

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- **The R help system**
- Web resources for R
- IDE editors for R


3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

The HTML help system

R has a comprehensive HTML help facility

- Run the `help.start` function
- R GUI menu item
Help|Html help



R Code: Starting HTML help

```
> help.start()
```

If nothing happens, you should open
`'http://127.0.0.1:12534/doc/html/index.html'` yourself

The help function

Obtain help on a particular topic via the help function

- `help(topic)`
- `?topic`

R Code: Topic help

```
> help(read.table)
```

The screenshot shows the R help page for the 'read.table' function in Mozilla Firefox. The URL is <http://127.0.0.1:17332/library/utils/html/read.table.html>. The page title is 'R: Data Input - Mozilla Firefox'. The content includes the function signature, detailed descriptions of parameters like 'file', 'header', and 'sep', and examples of how to use the function.

```
## R: Data Input - Mozilla Firefox
@... http://127.0.0.1:17332/library/utils/html/read.table.html
@... Google
@... R: Data Input
@... Documentation

Data Input

Description
Reads a file in table format and creates a data frame from it, with cases corresponding to lines and variables to fields in the file.

Usage
read.table(file, header = FALSE, sep = "", quote = "\\"", 
          dec = ".", row.names, col.names,
          na.la = "stringsAsFactors",
          na.strings = NA, comment.char = "#",
          nmax = 1000, check.names = TRUE, fill = blank.lines.skip,
          strip.white = TRUE, blank.lines.skip = TRUE,
          fill = FALSE),
          na.strings = NA, quote = "\\"", dec = ".",
          fill = TRUE, blank.lines.skip = TRUE),
          fileencoding = "", encoding = "unknown"),
          fill = TRUE, comment.char = "#")
read.csv(file, header = TRUE, sep = "\t", quote = "\"\"", dec = ",",
        fill = TRUE, comment.char = "#")
read.csv2(file, header = TRUE, sep = "\t", quote = "\"\"", dec = ",",
         fill = TRUE, comment.char = "#")
read.delim(file, header = TRUE, sep = "\t", quote = "\"\"", dec = ",",
           fill = TRUE, comment.char = "#")
read.delim2(file, header = TRUE, sep = "\t", quote = "\"\"", dec = ",",
            fill = TRUE, comment.char = "#")

Arguments
file
the name of the file which the data are to be read from. Each row of the table appears as one line of the file. If it does not contain an absolute path, the file name is relative to the current working directory. ?file\(\). Tilde-expansion is performed where supported. As from R 2.2.0 this can be a compressed file (?file\(\)).
Alternatively, file can be a readable text mode connection (which will be opened for reading if necessary), and if no close=TRUE (and hence destroyed) at the end of the function call (?close\(\)) is used, the prompt for lines may be necessary causing terminate input with a blank line or an EOF signal. close=TRUE is Unix and close=0 on Windows. Any pushback on readLines will be cleared before return.
fill
can also be a complete URL.
header
a logical value indicating whether the file contains the names of the variables as its first line. If missing, the value is determined from the file format. header is set to TRUE if and only if the first row contains one fewer field than the number of columns.
sep
the field separator character. Values on each line of the file are separated by this character. If sep = " (the default for read.csv) the separator is "white space", that is one or more spaces, tabs, newlines or carriage returns.
Done
```


The help.search function

Search help for a particular topic via the `help.search` function

- `help.search(topic)`
- `??topic`

R Code: Search help

```
> ??predict
```



A screenshot of a Windows desktop showing the R Help search interface. The title bar says "R Help search results". The main area is titled "Search Results" with a large R logo icon. Below it, a message says "The search string was \"predict\"". Under "Code demonstrations:", there is a link to "stats::glm.vr" which describes some `glm` examples from V&R with several predictors. There is also a link "(Run demo in console)". Under "Help pages:", there is a list of functions related to prediction:

Function	Description
<code>cluster::predict.ellipoid</code>	Predict Method for Ellipsoid Objects
<code>GARCH::predict-methods</code>	GARCH Prediction Function
<code>forecast::dm.test</code>	Diebold-Mariano test for predictive accuracy
<code>MASS::dose.p</code>	Product Doses for Binomial Assay model
<code>MASS::predict.glmPQL</code>	Predict Method for <code>glmPQL</code> Fits
<code>MASS::predict.lda</code>	Classify Multivariate Observations by Linear Discrimination
<code>MASS::predict.lpm</code>	Predict from an lpm Fit
<code>MASS::predict.mle</code>	Predict Method for Class 'mle'
<code>MASS::predict.qda</code>	Classify from Quadratic Discriminant Analysis
<code>MASS::rlm</code>	Robust Fitting of Linear Models

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- **Web resources for R**
- IDE editors for R


3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

R Homepage

<http://www.r-project.org>

- List of CRAN mirror sites
- Manuals
- FAQs
- Mailing Lists
- Links



CRAN - Comprehensive R Archive Network

<http://cran.fhcrc.org>

- CRAN Mirrors

- About 88 sites worldwide
- About 19 sites in US

- R Binaries

- R Packages

- R Sources

- Task Views

The screenshot shows the CRAN homepage (<http://cran.fhcrc.org>) in Mozilla Firefox. The page title is "The Comprehensive R Archive Network". On the left, there's a sidebar with links to "cran", "Mirrors", "What's new?", "Task Views", "Search", "About R", "R Homepage", and "The R Journal". Below that are links for "Software", "R Sources", "R Binaries", "Packages", "Other", "Documentation", "Manuals", "FAQs", and "Consultation". The main content area has a heading "Frequently used pages" with a link to "Download and Install R". It explains that precompiled binary distributions of the base system and contributed packages are available for Windows and Mac users. It lists links for "Linux", "MacOSX", and "Windows". A section titled "Source code for all platforms" discusses how Windows and Mac users can get source code that needs compilation. It lists links for "The latest release (2010-05-31): R-2.11.1.tar.gz" (with a note about what's new), "R alpha and beta releases" (daily snapshots), "Daily snapshots of current patched and development versions" (with a note about new features and bug fixes), "Source code of older versions of R" (available here), and "Contributed extension packages". A "Questions About R" section includes a link to "What are R and CRAN?". At the bottom, it provides a brief overview of R as "GNU S", stating it's a freely available language and environment for statistical computing and graphics which provides a wide variety of statistical and graphical techniques: linear and nonlinear modelling, statistical tests, time series analysis, classification, clustering, etc. It encourages users to consult the [R project homepage](#) for further information.

CRAN Task Views

Organizes the 3500+ R packages by application

- Finance
- Time Series
- Econometrics
- Optimization
- Machine Learning

The screenshot shows a Mozilla Firefox browser window displaying the CRAN Task Views page. The URL in the address bar is <http://cran.r-project.org>. The page title is "The Comprehensive R Archive Network - Mozilla Firefox". The main content area is titled "CRAN Task Views". On the left, there is a sidebar with links to various sections: "cran", "Mirrors", "What's new?", "Task Views", "Search", "About R", "R Homepage", "The R Journal", "Software", "R Sources", "R Binaries", "Packages", "Other", "Documentation", "Manuals", "FAQs", and "Contacted". The main content area lists task views categorized under "Bayesian", "Chemosometrics", "Clinical Trials", "Cluster", "Distributions", "Econometrics", "Environmetrics", "ExperimentalDesign", "Finance", "Genetics", "Graphics", "gl", "HighPerformanceComputing", "MachineLearning", "MedicalImaging", "Multivariate", "NaturalLanguageProcessing", "Optimization", "Pharmacokinetics", "Phylogenetics", "Psychometrics", "Robust", "SocialSciences", "Spatial", "Survival", and "TimeSeries". To the right of the categories, their respective descriptions are listed. At the bottom of the page, there is a block of text providing instructions for automatically installing the views.

```
To automatically install these views, the crv package needs to be installed, e.g., via  
install.packages("crv")  
library(crv)  
and then the views can be installed via install.views or update.views (which first assesses which of the  
packages are already installed and up-to-date), e.g.,  
install.views("Econometrics")  
or  
update.views("Econometrics")
```

Stackoverflow

Stackoverflow has become the primary resource for help with R

The screenshot shows a Mozilla Firefox browser window with the title bar "Recent Y Questions - Stack Overflow - Mozilla Firefox". The address bar contains "stackoverflow.com/questions/tagged/r". The page itself is the StackExchange homepage for R, featuring the StackOverflow logo and navigation links for "Questions", "Tags", "Users", "Badges", "Unanswered", and "Ask Question". A large counter on the right indicates "13,161 questions tagged". Below the counter, a link leads to "about".

The main content area displays three search results:

- pass value from R (run in SPSS) to open SPSS dataset**
0 votes, 0 answers, 6 views. Last updated 5 mins ago by DocBuckets.
- How to Apply lapply() or lapply() to component of list produced by lapply() or lapply()**
0 votes, 0 answers, 13 views. Last updated 50 mins ago by Kirk Fleming.
- load data set automatically**
0 votes, 0 answers, 20 views. Last updated 1 hour ago by Tyler Rinker.

On the right side, there is a "Community Bulletin" section with news items about the 2012 Community Moderator Election and the start of the 2012 Stack Overflow Community Moderator Election. Below that is a "Related Tags" sidebar listing various R-related tags with their counts: ggplot2 (1255), data.frame (711), plot (649), statistics (450), matrix (261), time-series (249), plyr (231), python (206), and r (187).

<http://stackoverflow.com/>

R-SIG-FINANCE

- Nerve center of the R finance community
- Daily must read
- Exclusively for Finance-specific questions, not general R questions

The screenshot shows a Mozilla Firefox browser window displaying the R-SIG-Finance Info Page. The URL in the address bar is <https://stat.ethz.ch/mailman/listinfo/r-sig-finance>. The page title is "R-SIG-Finance Info Page". The main content area has a blue header bar with the text "R-SIG-Finance -- Special Interest Group for 'R in Finance'". Below this, there's a section titled "About R-SIG-Finance" and another titled "Using R-SIG-Finance". A note says "To see the collection of prior postings to the list, visit the [R-SIG-Finance Archives](#)". There's also a note about sending messages to the list. A "Subscribe to R-SIG-Finance" section contains fields for email address, name (optional), and password, along with checkboxes for language preference and daily digest options. A "Subscribe" button is at the bottom.

<https://stat.ethz.ch/mailman/listinfo/r-sig-finance>

Google's R Style Guide

<http://google-styleguide.googlecode.com/svn/trunk/google-r-style.html>

- Naming convention
- Coding Syntax
- Program Organization

The screenshot shows the Google's R Style Guide page. At the top, there is a navigation bar with links for Home, Style Guide, and Help. Below the navigation, the title "Google's R Style Guide" is displayed, followed by a brief description: "R is a high-level programming language used primarily for statistical computing and graphics. The goal of the R Programming Style Guide is to make our R code easier to read, share, and verify. The rules below were designed in collaboration with the entire R team community at Google." A "Summary: R Style Rules" section follows, containing 14 numbered items from 1 to 14, each with a small icon and a brief description. Below this is another summary section for "R Language Rules" with 3 numbered items. At the bottom, there are two sections: "1. Notation and Naming" which includes "File Names" and "Identifiers", and "2. Functionality" which includes "Functions", "Methods", and "Imports". Each section contains examples of good and bad practices.

Summary: R Style Rules

1. **File Names**: avoid using .R
2. **Identifiers**: variables, name, FunctionName, wConsistentName
3. **Line Length**: maximum 80 characters
4. **Indentation**: two spaces, no tabs
5. **Semicolons**: ;
6. **Group Braces**: first on same line, last on own line
7. **Assignment**: use <-, not =
8. **Separators**: don't use ;
9. **General Layout and Ordering**
10. **Blank Lines**: all comments begin with a space; idea comments need two spaces before the #
11. **Function Definitions and Calls**
12. **Function Documentation**
13. **Example Facilities**
14. **TODO Style**: TODO (comment)

Summary: R Language Rules

1. **return**: avoid using it
2. **functions**: errors should be raised using stop()
3. **Objects and Methods**: avoid S4 objects and methods when possible; never mix S3 and S4

1. Notation and Naming

File Names

File names should end in .R and, of course, be meaningful.
GOOD: `predict_ad_covariance.R`
BAD: `foo.R`

Identifiers

Don't use underscores (_) or hyphens (-) in identifiers. Identifiers should be named according to the following conventions. Variable names should have all lower case letters and words separated with dots (.); function names have initial capital letters and no dots (CapWords); constants are named like the function but with all initial caps.

- **variable_name**
GOOD: `avg_clicks`
BAD: `avg_clicks_, avg_clicks`
- **FunctionName**
GOOD: `calculateAvgClicks`

Quick R

<http://www.statmethods.net>

Introductory R Lessons

- R Interface
- Data Input
- Data Management
- Basic Statistics
- Advanced Statistics
- Basic Graphs
- Advanced Graphs

The screenshot shows the Quick-R website (http://www.statmethods.net) in a Mozilla Firefox browser window. The page has a dark green header with the title "Quick-R" and a subheader "for SAS/SPSS/Stata Users". Below the header is a navigation bar with links: Home, Interface, Input, Manage, Stats, Adv Stats, Graphs, and Adv Graphs. The main content area features several sections: "About Quick-R" (with a small R logo icon), "Why Use R?", "Obtaining R", "Quick-R as a book" (with a small book icon), and "Feedback". A sidebar on the right is titled "Top Menu" and lists various statistical topics. Another sidebar titled "Search" contains a search input field. At the bottom, there's a footer with copyright information and links to "Home" and "Sitemap".

R graphics details, colors, and other tech notes

R Graphics and other useful information by Earl Glynn of Stowers Institute for Medical Research

- URL

<http://research.stowers-institute.org/efg/R/index.htm>

- Features

- R Color Chart
- Using Color in R (great presentation)
- Plot area, margins, multiple figures
- Mixture models
- Distance measures and clustering
- Using Windows Explorer to Start R with Specified Working Directory (under tech notes)

Programming in R

Online R programming manual from UC Riverside

- URL

<http://manuals.bioinformatics.ucr.edu/home/programming-in-r>

- Selected Topics

- R Basics
- Finding Help
- Code Editors for R
- Control Structures
- Functions
- Object Oriented Programming
- Building R Packages

<http://rcom.univie.ac.at>

COM interface for R connectivity

- Excel
- Word
- C#
- VB
- Delphi

Download site for RAndFriends

- R
- Statconn
- Notepad++



Other useful R sites

R Bloggers	Aggregation of about 290 R blogs <ul style="list-style-type: none">● http://www.r-bloggers.com
R Site Search	Search R function help, vignettes, R-help <ul style="list-style-type: none">● http://finzi.psych.upenn.edu/search.html
R Seek	R specific search site <ul style="list-style-type: none">● http://www.rseek.org/
R Graph Gallery	Examples of many possible R graphs <ul style="list-style-type: none">● http://addictedtor.free.fr/graphiques
Revolution Blog	Blog from David Smith of Revolution <ul style="list-style-type: none">● http://blog.revolutionanalytics.com
Inside-R	R community site by Revolution Analytics <ul style="list-style-type: none">● http://www.inside-r.org

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R


3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

RStudio

RStudio is a fully-featured open-source IDE
for R

- R language highlighting
- Paste/Source code to R
- object explorer
- graphics window in main IDE




RStudio also provides a server-based version (R running in the cloud)

Revolution R Enterprize Visual Development Environment

Revolution Analytics is a company that sells a commercial distribution of R including a desktop IDE

Revolution R Enterprize is *free* to academic users

- R language highlighting
- Paste/Source code to R
- object explorer
- runs R in SDI mode




<http://www.revolutionanalytics.com>

WinEdt and R-Sweave

Based on WinEdt, an excellent shareware editor with support for L^AT_EX and Sweave development

- R language highlighting
- Paste/Source code to R
- Supports R in MDI mode
- Paste/Source code to S-PLUS



The screenshot shows the WinEdt interface with an R-Sweave document open. The code includes R library imports, a source command for a Sweave file, and various R assignments and functions. The interface features a menu bar, toolbars, and status bars at the bottom.

```
library(gbm)
library(TTR)
library(blotter)
library(PerformanceAnalytics)
library(zoo)
library(xts)
library(dyn)
library(gbm)

#m.list<-list()
#try(m.list<-ls(envir=.blotter),envir=.blotter),silent=TRUE)
#try(m.list<-ls(envir=.instrument),envir=.instrument),silent=TRUE

# residualCointTest.r
source("residualCointTest.r")

graphics.off()
windows(height=8,width=10)

# constants
starting.date = '2002-01-10'
ending.date = '2009-12-31'

kInitialEquity = 1e7
kTradeSize = kInitialEquity/10
kWindowSize = 60
kTnnFee = 0.005

kExtremeHigh = 0.90
kThresholdHigh = 0.50
kThresholdLow = 0.50
kExtremeLow = 0.10
kMinStat = -1.5
kMinFV = 1.0

# setup account, portfolio, and pair, and pair, and data
pair.list = list(c("LLTC","ADI"),c("XLNK","ADI"),c("ALTR","ALTR"),c("X"))
pair.vec = unlist(lapply(X=pair.list, FUN = function(x) paste(x,collapse="")))
```

<http://www.winedt.com>


<http://www.winedt.org/Config/modes/R-Sweave.php>

StatET - An Eclipse Plug-In for R

StatET is a plug-in for the open-source Eclipse development environment

- R language highlighting
- Paste/Source code to R
- Source code debugger
- Supports R in SDI mode
- Excellent documentation by Longhow Lam


<http://www.walware.de/goto/statet>



Notepad++ and NpptoR

NpptoR is an automation widget
(based on AutoHotkey) which allows
the very useful program editor
Notepad++ to interact with R

- R language highlighting
- Paste/Source code to R
- Supports R in SDI mode
- Can be installed as part of RAndFriends



A screenshot of the Notepad++ application window. The title bar reads "C:\Rprojects\UNW\download\Yahoo\script01.R - Notepad++". The menu bar includes File, Edit, Search, View, Encoding, Language, Settings, Macro, Run, TexFx, Plugins, Window, and Help. The toolbar has icons for Open, Save, Find, Replace, Cut, Copy, Paste, Select All, and others. The status bar at the bottom shows "R programming language | 597 chars | 631 bytes | 18 lines | Ln: 12 Col: 1 Sel: 0 (0 bytes) in 0 ranges | Dos/Windows ANSI | INS". The main text area contains R code:

```
1 setwd(paste("D:\\UN\\VCA",sep="\\"))
2
3 imb <- get.hist.quote(instrument="IBM",quote="A",start="2010-01-01",quiet=TRUE)
4 hpq <- get.hist.quote(instrument="HPQ",quote="A",start="2010-01-01",quiet=TRUE)
5 msft <- get.hist.quote(instrument="MSFT",quote="A",start="2010-01-01",quiet=TRUE)
6 intc <- get.hist.quote(instrument="INTC",quote="A",start="2010-01-01",quiet=TRUE)
7
8 dat <- merge(imb,hpq,msft,intc)
9 colnames(dat) <- c("IBM","HPQ","MSFT","INTC")
10 head(dat)
11 write.zoo(x=dat,file="ilmhpq.txt")
12
13 rm(dat)
14
15 prices <- read.zoo(file="ilmhpq.txt",header=TRUE)
16 returns <- 100*diff(log(prices))
17 pca.mod <- prcomp(returns,center=TRUE,scale.= FALSE)
```

<http://notepad-plus-plus.org>

<http://sourceforge.net/projects/npptor>

<http://rcom.univie.ac.at/download.html>

Other R IDE links

Tinn-R Popular R IDE

- <http://www.sciviews.org/Tinn-R>

ESS Emacs Speaks Statistics

- <http://ess.r-project.org>

other R GUI Projects

- http://www.sciviews.org/_rgui

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- **Basic plotting**
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

Basic plotting functions

Function	Description
plot	generic function to plot an R object
lines	adds lines to the current plot
segments	adds lines line segments between point pairs
points	adds points to the current plot
text	adds text to the current plot
abline	adds straight lines to the current plot
curve	plot a function over a range
legend	adds a legend to the current plot
matplot	plot all columns of a matrix
par	sets graphics parameters

The plot function

The plot function is a generic function for plotting of R objects

R Code: plot arguments

```
> args(plot.default)

function (x, y = NULL, type = "p", xlim = NULL, ylim = NULL,
  log = "", main = NULL, sub = NULL, xlab = NULL, ylab = NULL,
  ann = par("ann"), axes = TRUE, frame.plot = axes, panel.first = NULL,
  panel.last = NULL, asp = NA, ...)
NULL
```

x vector to be plotted (or index if y given)

y vector to be plotted

xlim/ylim x & y limited

xlab/ylab x & y axis labels

main plot title (can be done with title function)

type "p" = points (default), "l" = lines, "h" = bars, "n" = no plot


col color or bars

asp control the aspect ratio

The plot function

R Code: Plot with defaults


```
> library(Ecdat)  
> data(Capm)  
> plot(Capm[, "rf"])
```



The plot function

R Code: Plot lines


```
> plot(Capm[, "rf"], type="l")
```



The plot function

R Code: Plot bars


```
> plot(Capm[, "rmrf"], type="h")
```



The plot function

R Code: XY plot

```
> plot(Capm[, "rmrf"] , Capm[, "rcon"])
```



The points function

The points function adds points to the current plot at the given x, y coordinates

R Code: points arguments

```
> args(points.default)  
function (x, y = NULL, type = "p", ...)  
NULL
```

x vector of x coordinates

y vector of y coordinates

The lines function

The `lines` function adds connected line segments to the current plot

R Code: `lines` arguments

```
> args(lines.default)  
  
function (x, y = NULL, type = "l", ...)  
NULL
```

- x** vector of x coordinates
- y** vector of y coordinates

The text function

The `text` function adds text labels to a plot at given `x`, `y` coordinates

R Code: `text` arguments

```
> args(text.default)

function (x, y = NULL, labels = seq_along(x), adj = NULL, pos = NULL,
  offset = 0.5, vfont = NULL, cex = 1, col = NULL, font = NULL,
  ...)
NULL
```

`x/y` location to place text

`labels` text to be display

`adj` adjustment of label at `x`, `y` location


`pos` position of text relative to `x`, `y`

`offset` offset from `pos`

Plotting a blank frame

R Code: Plotting a blank frame


```
> plot(0,xlim=c(-20,20),ylim=c(-20,20),type="n",
      xlab="market return",ylab="construction return")
```



A blank frame with points, lines, and text added

R Code: Adding points, lines, and text to a blank frame

```
> plot(0,xlim=c(-20,20),ylim=c(-20,20),type="n",
+     xlab="market return",ylab="construction return")
> points(x=Capm[, "rmrf"],y=Capm[, "rcon"],col="gray")
> lines(x=-20:20,y=-20:20,lwd=2,col="darkred")
> text(20,20,labels="slope = 1",pos=2)
```



The segments function

The segments function draws line segments between point pairs

R Code: segments arguments

```
> args(segments)

function (x0, y0, x1 = x0, y1 = y0, col = par("fg"), lty = par("lty"),
  lwd = par("lwd"), ...)
NULL
```

`x0, y0` point coordinates from which to draw

`x1, y1` point coordinates to which to draw

The curve function

The curve function draws a curve of a function or expression over a range

R Code: curve arguments

```
> args(curve)

function (expr, from = NULL, to = NULL, n = 101, add = FALSE,
  type = "l", xname = "x", xlab = xname, ylab = NULL, log = NULL,
  xlim = NULL, ...)
NULL
```

expr function or expression of x

from start of range

to end of range

n number of points over from/to range

add add to current plot (T/F)

The abline function

The abline function adds one or more straight lines through the current plot

R Code: abline arguments

```
> args(abline)  
  
function (a = NULL, b = NULL, h = NULL, v = NULL, reg = NULL,  
        coef = NULL, untf = FALSE, ...)  
NULL
```

h/v vertical or horizontal coordinate of line

a/b intercept and slope of line

The `matplot` function

The `matplot` function plots multiple columns of a matrix versus an index

R Code: `matplot` arguments

```
> args(matplot)

function (x, y, type = "p", lty = 1:5, lwd = 1, lend = par("lend"),
  pch = NULL, col = 1:6, cex = NULL, bg = NA, xlab = NULL,
  ylab = NULL, xlim = NULL, ylim = NULL, ..., add = FALSE,
  verbose = getOption("verbose"))
NULL
```

`x/y` matrices or vectors to be plotted

Graphical parameters controlled via the par function

R is capable of producing publication quality graphics by allowing (*requiring*) fine-grained control of a number of graphics parameters

R Code: Names of graphical parameters

```
> names(par())
```

```
[1] "xlog"      "ylog"       "adj"        "ann"        "ask"        "bg"
[7] "bty"        "cex"         "cex.axis"    "cex.lab"    "cex.main"    "cex.sub"
[13] "cin"        "col"         "col.axis"    "col.lab"    "col.main"    "col.sub"
[19] "cra"        "crt"         "csi"         "cxy"        "din"        "err"
[25] "family"     "fg"          "fig"         "fin"        "font"       "font.axis"
[31] "font.lab"   "font.main"   "font.sub"    "lab"        "las"        "lend"
[37] "lheight"    "ljoin"       "lmitre"     "lty"        "lwd"        "mai"
[43] "mar"        "mex"         "mfcol"      "mfg"        "mfrow"      "mpg"
[49] "mkh"        "new"         "oma"        "omd"        "omi"        "pch"
[55] "pin"        "plt"         "ps"         "pty"        "smo"        "srt"
[61] "tck"        "tcl"         "usr"        "xaxp"      "xaxs"      "xaxt"
[67] "xpd"        "yaxp"       "yaxs"      "yaxt"      "ylbias"
```

Commonly used par parameters

Parameter	Description
col	plot color
lwd	line width
lyt	line type
mfrow	set/reset multi-plot layout
cex.axis	character expansion - axis
cex.lab	character expansion - labels
cex.main	character expansion - main
pch	point character
las	axis label orientation
bty	box type around plot or legend

- some parameters can be passed in a plot function (e.g. col, lwd)
- some parameters can only be changed by a call to par (e.g. mfrow)

The legend function

R Code: legend arguments

```
> args(legend)

function (x, y = NULL, legend, fill = NULL, col = par("col"),
  border = "black", lty, lwd, pch, angle = 45, density = NULL,
  bty = "o", bg = par("bg"), box.lwd = par("lwd"), box.lty = par("lty"),
  box.col = par("fg"), pt.bg = NA, cex = 1, pt.cex = cex, pt.lwd = lwd,
  xjust = 0, yjust = 1, x.intersp = 1, y.intersp = 1, adj = c(0,
    0.5), text.width = NULL, text.col = par("col"), text.font = NULL,
  merge = do.lines && has.pch, trace = FALSE, plot = TRUE,
  ncol = 1, horiz = FALSE, title = NULL, inset = 0, xpd, title.col = text.col,
  title.adj = 0.5, seg.len = 2)
NULL
```

- x/y** location of the legend (can be give as a position name)
- legend** vector of labels for the legend
- col** vector of colors
- lty** line type
- lwd** line width
- pch** character

The barplot function

The barplot function can create vertical or horizontal barplots

R Code: barplot arguments

```
> args(barplot.default)

function (height, width = 1, space = NULL, names.arg = NULL,
  legend.text = NULL, beside = FALSE, horiz = FALSE, density = NULL,
  angle = 45, col = NULL, border = par("fg"), main = NULL,
  sub = NULL, xlab = NULL, ylab = NULL, xlim = NULL, ylim = NULL,
  xpd = TRUE, log = "", axes = TRUE, axisnames = TRUE, cex.axis = par("cex.axis")
  cex.names = par("cex.axis"), inside = TRUE, plot = TRUE,
  axis.lty = 0, offset = 0, add = FALSE, args.legend = NULL,
  ...)
NULL
```

height vector or matrix (stacked bars or side-by-side bars) of heights

names.arg axis labels for the bars

beside stacked bars or side-by-side if height is a matrix

legend vector of labels for stacked or side-by-side bars

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- **Basic statistics and the normal distribution**
- Working with time series in R
- Variable scoping in R

Probability distributions

- Random variable

A *random variable* is a quantity that can take on any of a set of possible values but only one of those values will actually occur

- *discrete* random variables have a finite number of possible values
- *continuous* random variables have an infinite number of possible values

- Probability distribution


The set of all possible values of a random variable along with their associated probabilities constitutes a *probability distribution* of the random variable

PDFs and CDFs

- Probability density function (PDF)


$$Pr(a < Y < b) = \int_a^b f_Y(y) dy$$

$$\int_{-\infty}^{\infty} f_Y(y) dy = 1$$






- Cumulative distribution function (CDF)

$$F_Y(y) = Pr(Y \leq y) = \int_{-\infty}^y f_Y(y) dy$$



PDF, CDF, quantile functions

Function	General Notation	Normal Notation	R	Excel	Graph
pdf	$f(x)$	$\phi(z)$	dnorm	NORMDIST	
cdf	$F(x)$	$\Phi(z)$	pnorm	NORMDIST	
quantile	$F^{-1}(x)$	$\Phi^{-1}(z)$	qnorm	NORMINV	

Normal distribution PDF function: dnorm

dnorm computes the normal PDF: $\phi(z)$

R Code: Plot PDF


```
> args(dnorm)
function (x, mean = 0, sd = 1, log = FALSE)
NULL

> x <- seq(from = -5, to = 5, by = 0.01)
> x[1:10]
[1] -5.00 -4.99 -4.98 -4.97 -4.96 -4.95 -4.94 -4.93 -4.92 -4.91

> y <- dnorm(x)
> y[1:5]
[1] 1.486720e-06 1.562867e-06 1.642751e-06 1.726545e-06 1.814431e-06

> par(mar = par()$mar + c(0,1,0,0))
> plot(x=x,y=y,type="l",col="seagreen",lwd=2,
  xlab="x",ylab="density\ny = dnorm(x)")
> grid(col="darkgrey",lwd=2)
> title(main="Probability Density Function (PDF)")
```

Normal distribution PDF function: dnorm



Others:

- dt
- dstd
- dsstd
- dged
- dsged
- dst
- dmst
- dct

Normal distribution CDF functions: `pnorm` and `qnorm`

`pnorm` computes the normal CDF:

$$\Pr(X \leq z) = \Phi(z)$$

`qnorm` computes the inverse of the normal CDF (i.e. quantile):

$$z_\alpha = \Phi^{-1}(\alpha)$$


R Code: Plot CDF

```
> args(pnorm)
function (q, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)
NULL

> args(qnorm)
function (p, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)
NULL

> y <- pnorm(x)
> par(mar = par()$mar + c(0,1,0,0))
> plot(x=x,y=y,type="l",col="seagreen",lwd=2, xlab="x = qnorm(y)",
       ylab="probability\ny = pnorm(x)" ; grid(col="darkgrey",lwd=2)
> title(main="Cumulative Distribution Function (CDF)")
```

Normal distribution CDF functions: `pnorm` and `qnorm`



Generating normally distributed random numbers

The function `rnorm` generates random numbers from a normal distribution

R Code: `rnorm` arguments

```
> args(rnorm)

function (n, mean = 0, sd = 1)
NULL

> x <- rnorm(150)
> x[1:5]

[1] -0.6264538  0.1836433 -0.8356286  1.5952808  0.3295078

> y <- rnorm(50, sd=3)
> y[1:5]

[1]  1.3505613 -0.0556795 -0.9542051 -2.7880864 -4.4623809
```

`n` number of observations

`mean` mean of distribution

`sd` standard deviation of distribution

Histograms

The generic function `hist` computes a histogram of the given data values

R Code: `hist` arguments

```
> args(hist.default)

function (x, breaks = "Sturges", freq = NULL, probability = !freq,
  include.lowest = TRUE, right = TRUE, density = NULL, angle = 45,
  col = NULL, border = NULL, main = paste("Histogram of", xname),
  xlim = range(breaks), ylim = NULL, xlab = xname, ylab, axes = TRUE,
  plot = TRUE, labels = FALSE, nclass = NULL, warn.unused = TRUE,
  ...)
NULL
```

`x` vector of histogram data

`breaks` number of breaks, vector of breaks, name of break algorithm, break function

`prob` probability densities or counts


`ylim` y-axis range

`col` color or bars

Plotting histograms

R Code: Plotting histograms


```
> hist(x,col="seagreen")
```



Plotting histograms

R Code: Plotting histograms


```
> hist(c(x,y),prob=T,breaks="FD",col="seagreen")
```



Plotting histograms

R Code: Plotting histograms

```
> hist(c(x,y),prob=T,breaks=50,col="seagreen")
```



Basic stats functions

Short list of some common statistics and math functions:

- `mean` mean of a vector or matrix
- `median` median of a vector or matrix
- `mad` median absolute deviation of a vector or matrix
- `var` variance of a vector or matrix
- `sd` standard deviation of a vector
- `cov` covariance between vectors
- `cor` correlation between vectors
- `diff` difference between elements in a vector
- `log` log of a vector or matrix
- `exp` exponentiation of a vector or matrix
- `abs` absolute value of a vector or matrix

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- **Working with time series in R**
- Variable scoping in R

Time series data

Time Series

A time series is a sequence of *ordered* data points measured at specific points in time

A time series class in R is a *compound data object* that includes a data matrix as well as a vector of associated time stamps

class	package	overview
ts	base	regularly spaced time series
mts	base	multiple regularly spaced time series
timeSeries	rmetrics	default for Rmetrics packages
zoo	zoo	reg/irreg and arbitrary time stamp classes
xts	xts	an extension of the zoo class

Time series methods

Time series classes in R will typically implement the following methods:

- start** return start of time series
- end** return end of time series
- frequency** return frequency of time series
- window** Extract subset of time series
- index** return time index of time series
- time** return time index of time series
- coredata** return data of time series
- diff** difference of the time series
- lag** lag of the time series
- aggregate** aggregate to lower resolution time series
- cbind** merge 2 or more time series together

Creating a zoo object

R Code: Creating a zoo object

```
> library(zoo)
> msft.df <- read.table("MSFT.CSV", header = TRUE, sep = ",", as.is = TRUE)
> head(msft.df,2)

      Date  Open  High  Low Close   Volume Adj.Close
1 2009-09-21 25.11 25.37 25.1 25.30 28864500     25.30
2 2009-09-18 25.46 25.48 25.1 25.26 68016500     25.26

> args(zoo)

function (x = NULL, order.by = index(x), frequency = NULL)
NULL

> msft.z <- zoo(x=msft.df[, "Close"],order.by=as.Date(msft.df[, "Date"]))
> head(msft.z)

2008-09-22 2008-09-23 2008-09-24 2008-09-25 2008-09-26 2008-09-29
      25.40      25.44      25.72      26.61      27.40      25.01
```

Inspecting a zoo object


R Code: Inspecting a zoo object

```
> class(msft.z)
[1] "zoo"
> start(msft.z)
[1] "2008-09-22"
> end(msft.z)
[1] "2009-09-21"
> frequency(msft.z)
[1] 1
> class(coredata(msft.z))
[1] "numeric"
> class(time(msft.z))
[1] "Date"
```

Plotting a zoo object

R Code: Plotting a zoo object

```
> ticks <- as.Date(unique(as.yearmon(index(msft.z))))  
> plot(msft.z, xaxt='n', xlab="", ylab="$", main="Microsoft Stock Price")  
> axis(side=1, at=ticks, lab=format(ticks,"%b-%y"))
```



The read.zoo function

R Code: reading a file into a zoo object

```
> args(read.zoo)

function (file, format = "", tz = "", FUN = NULL, regular = FALSE,
  index.column = 1, drop = TRUE, FUN2 = NULL, split = NULL,
  aggregate = FALSE, ..., text)
NULL

> soft <- read.zoo(file="MSFT.CSV",header=TRUE,sep=",")
> head(soft,2)

      Open   High    Low Close   Volume Adj.Close
2008-09-22 26.22 26.32 25.32 25.40 105207700     24.76
2008-09-23 25.66 26.17 25.34 25.44  92181300     24.80

> class(soft)

[1] "zoo"

> class(coredata(soft))

[1] "matrix"

> class(index(soft))

[1] "Date"
```

Outline

1 Part 1

- R overview and history
- R language references

2 Part 2

- R language and environment basics
- Data structures, data manipulation, working directory, data files
- The R help system
- Web resources for R
- IDE editors for R

3 Part 3

- Basic plotting
- Basic statistics and the normal distribution
- Working with time series in R
- Variable scoping in R

Free variables

In the body of a function, 3 types of symbols may be found:

- formal parameters - arguments passed in the function call
- local variables - variables created in the function
- free variables - variables created outside of the function
(note, free variables become local variables if you assign to them)

R Code: Types of variables in functions

```
> f <- function(x) {  
  y <- 2*x  
  print(x) # formal parameter  
  print(y) # local variable  
  print(z) # free variable  
}
```

Environments

The main workspace in R (i.e. what you are interacting with at the R console) is called the *global environment*

According to the scoping rules of R (referred to as *lexical scoping*), R will search for a free variable in the following order:

- ① The environment in which the function was created
 - For functions created in the global environment, this will be the global environment
- ② The parent environment of the environment where the function was created
- ③ The parent of the parent ... up until the global environment is searched
- ④ The search path of loaded libraries found using the `search()` function

The search path

The function `search` returns a list of attached packages which will be searched in order (after the global environment) when trying to resolve a free variable

R Code: The search path

```
> search()  
  
[1] ".GlobalEnv"           "package:zoo"          "package:Ecdat"  
[4] "package:nutshell"     "package:stats"       "package:graphics"  
[7] "package:grDevices"    "package:datasets"   "package:utils"  
[10] "Autoloads"            "package:base"
```

Variable scoping examples

R Code: Variable scoping examples

```
> # example 1  
> a <- 10  
> x <- 5  
> f <- function (x) x + a  
> f(2)
```

```
[1] 12
```

```
> # example 2  
> f<- function (x)  
{  
  a<-5  
  g(x)  
}  
> g <- function(y) y + a  
> f(2)
```

```
[1] 12
```

Variable scoping examples

R Code: Variable scoping examples

```
> # example 3
> f <- function (x) {
  a<-5
  g <- function (y) y + a
  g(x)
}
> f(2)
[1] 7

> # example 4
> f <- function (x) {
  x + mean(rivers) # rivers is defined in the dataset package
}
> f(2)
[1] 593.1844
```

W COMPUTATIONAL FINANCE & RISK MANAGEMENT

UNIVERSITY *of* WASHINGTON
Department of Applied Mathematics

<http://depts.washington.edu/comppfin>